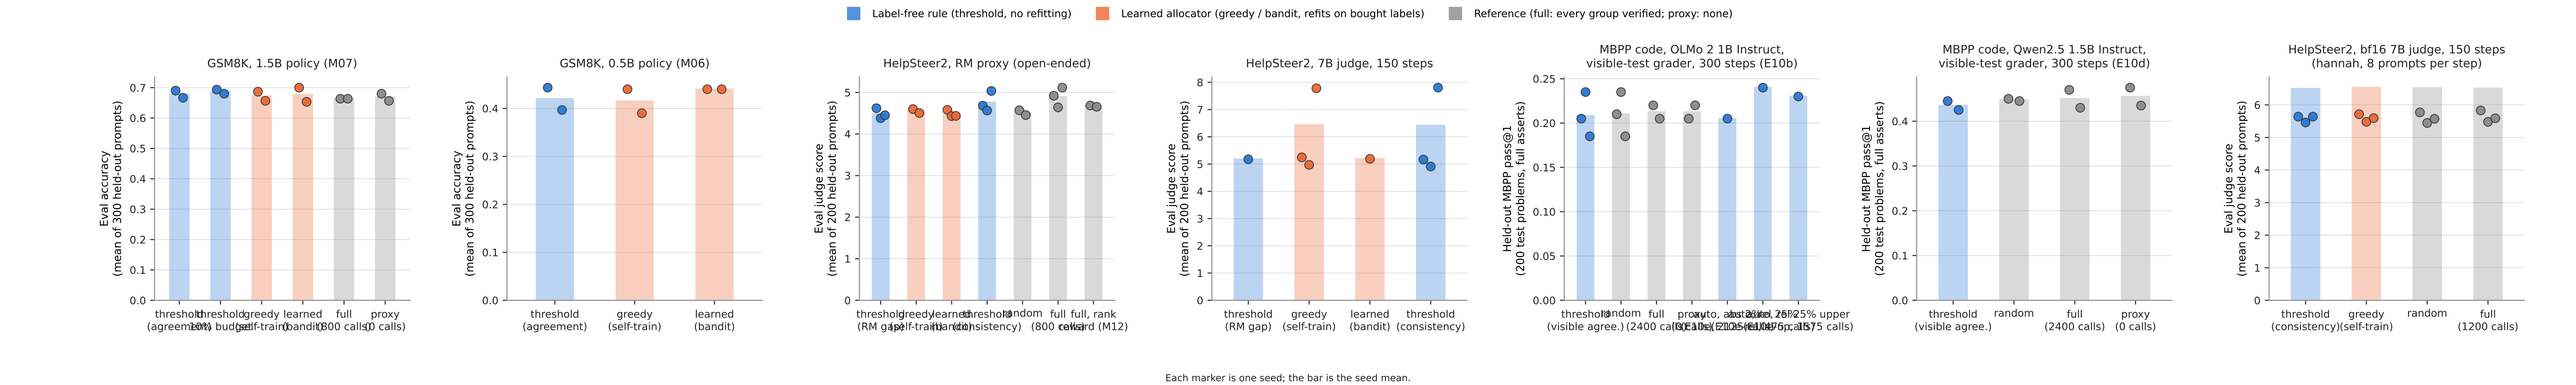
Label-free rule (threshold, no refitting)
Learned allocator (greedy / bandit, refits on bought labels)
Reference (full: every group verified; proxy: none)
GSM8K, 1.5B policy (M07)
Eval accuracy (mean of 300 held-out prompts)
threshold (agreement)
threshold (20% budget)
greedy (self-train)
learned (bandit)
full (300 calls)
proxy (0 calls)
GSM8K, 0.5B policy (M06)
Eval accuracy (mean of 300 held-out prompts)
threshold (agreement)
greedy (self-train)
learned (bandit)
HelpSteer2, RM proxy (open-ended)
Eval judge score (mean of 200 held-out prompts)
threshold (RM gap)
greedy (self-train)
learned (bandit)
threshold (consistency)
random
full (800 calls)
full, rank reward (M12)
HelpSteer2, 7B judge, 150 steps
Eval judge score (mean of 200 held-out prompts)
threshold (RM gap)
greedy (self-train)
learned (bandit)
threshold (consistency)
MBPP code, OLMo 2 1B Instruct, visible-test grader, 300 steps (E10b)
Held-out MBPP pass@1 (200 test problems, full asserts)
threshold (visible agree.)
random (2400 calls)
full (9610 calls)
proxy, abstain 25% (E10b, 2154 calls)
proxy, abstain 25% (E10b, 1875 calls)
proxy, abstain 25% upper (E10b, 1575 calls)
MBPP code, Qwen2.5 1.5B Instruct, visible-test grader, 300 steps (E10d)
Held-out MBPP pass@1 (200 test problems, full asserts)
threshold (visible agree.)
random (2400 calls)
full (9610 calls)
proxy (0 calls)
HelpSteer2, bf16 7B judge, 150 steps (hannah, 8 prompts per step)
Eval judge score (mean of 200 held-out prompts)
threshold (consistency)
greedy (self-train)
random (1200 calls)
full (4800 calls)
Each marker is one seed; the bar is the seed mean.