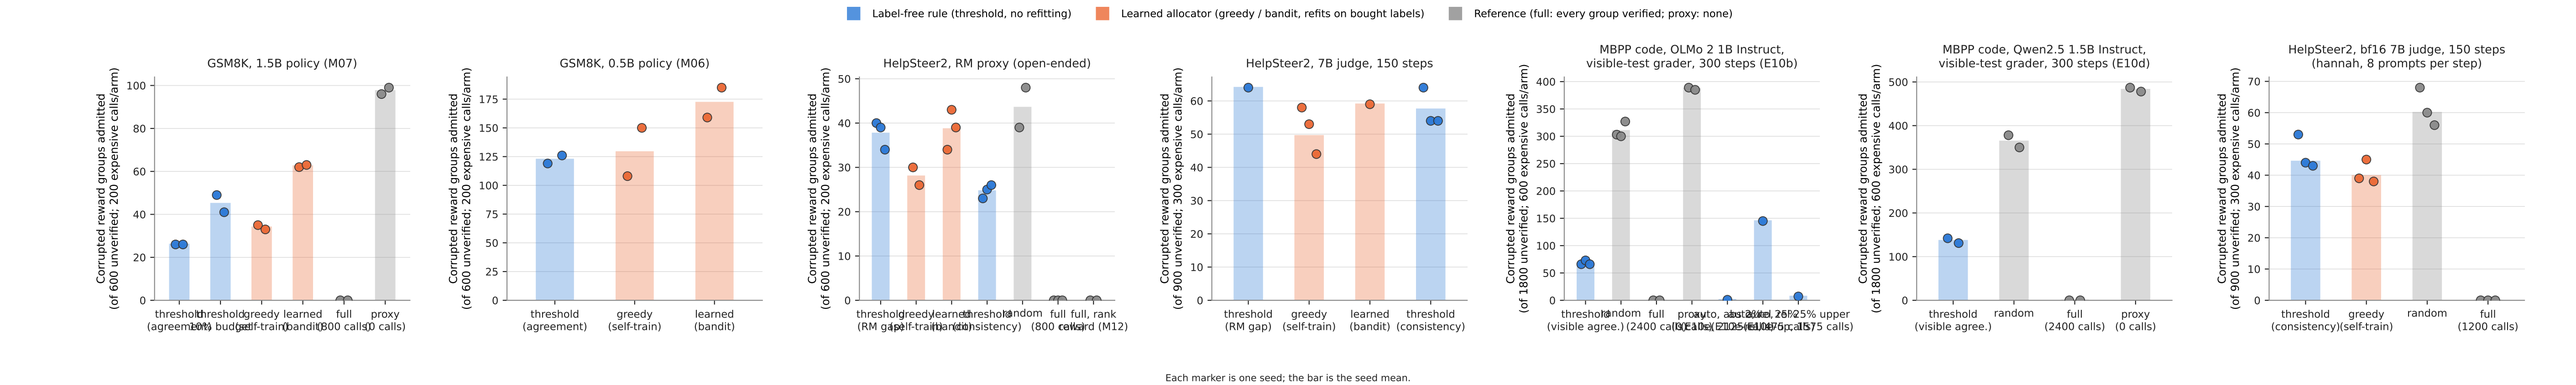
Label-free rule (threshold, no refitting)
Learned allocator (greedy / bandit, refits on bought labels)
Reference (full: every group verified; proxy: none)
GSM8K, 1.5B policy (M07)
Corrupted reward groups admitted (of 600 unverified; 200 expensive calls/arm)
threshold (agreement)
threshold (20% budget)
greedy (self-train)
learned (bandit)
full (600 calls)
proxy (0 calls)
GSM8K, 0.5B policy (M06)
Corrupted reward groups admitted (of 600 unverified; 200 expensive calls/arm)
threshold (agreement)
greedy (self-train)
learned (bandit)
HelpSteer2, RM proxy (open-ended)
Corrupted reward groups admitted (of 600 unverified; 200 expensive calls/arm)
threshold (RM gap)
greedy (self-train)
learned (bandit)
threshold (consistency)
random
full (800 calls)
full, rank reward (M12)
HelpSteer2, 7B judge, 150 steps
Corrupted reward groups admitted (of 900 unverified; 300 expensive calls/arm)
threshold (RM gap)
greedy (self-train)
learned (bandit)
threshold (consistency)
MBPP code, OLMo 2 1B Instruct, visible-test grader, 300 steps (E10b)
Corrupted reward groups admitted (of 1800 unverified; 600 expensive calls/arm)
threshold (visible agree.)
random
full (2400 calls)
proxy (0 calls)
proxy, abstain 25% (E10c)
learned, 25% (E10d)
upper (1575 calls)
MBPP code, Qwen2.5 1.5B Instruct, visible-test grader, 300 steps (E10d)
Corrupted reward groups admitted (of 1800 unverified; 600 expensive calls/arm)
threshold (visible agree.)
random
full (2400 calls)
proxy (0 calls)
HelpSteer2, bf16 7B judge, 150 steps (hannah, 8 prompts per step)
Corrupted reward groups admitted (of 900 unverified; 300 expensive calls/arm)
threshold (consistency)
greedy (self-train)
random
full (1200 calls)
Each marker is one seed; the bar is the seed mean.